

Code Running Guide

How to Run and Extend the Note 2 Python Examples

Installation, smoke tests, training diagnostics, outputs, and troubleshooting

Companion to the PINN/PIDL cyber-control lecture note

Contents

1	What This Guide Covers	2
2	Folder Structure	2
3	Installation	2
4	Quick Smoke Test	3
5	Normal Runs	3
5.1	Static figure generation	3
5.2	Inverse PINN	3
5.3	PIDL missing mechanism	3
5.4	Direct control PINN	3
5.5	PMP-informed PINN	3
5.6	Longer training diagnostics	4
6	How to Interpret Outputs	4
7	Output Map	4
8	Extending the Code	5
8.1	From sparse inverse PINN to noisy cyber data	5
8.2	From PIDL correction to richer missing mechanisms	5
8.3	From open-loop to feedback neural control	5
8.4	From small SIR models to network-scale PINNs	5
9	Troubleshooting	5
10	Reproducibility Checklist	5

1. What This Guide Covers

This guide is a standalone manual for the Python examples in Note 2. It explains how to set up the environment, run the fast checks, run longer PINN/PIDL diagnostics, interpret the generated outputs, and extend the examples to richer cyber-control models.

2. Folder Structure

The repository is organized around four teaching experiments:

```
.
├── requirements.txt
├── README.md
├── START_HERE.md
├── scripts/
│   ├── run_smoke_tests.sh
│   ├── generate_figures.py
│   └── run_training_iterations.py
├── src/
│   ├── inverse_pinn_sir_malware.py
│   ├── pidl_unknown_mechanism.py
│   ├── control_pinn_malware.py
│   └── pmp_informed_pinn_malware.py
├── experiments/
│   ├── OUTPUT_PREVIEW.md
│   ├── training_summary.md
│   ├── inverse_pinn_training_history.csv
│   ├── pidl_training_history.csv
│   ├── control_pinn_training_history.csv
│   └── pmp_informed_pinn_training_history.csv
├── figures/
│   ├── inverse_pinn_sparse_data.png
│   ├── pidl_missing_mechanism.png
│   ├── neural_architectures.png
│   └── training_iteration_diagnostics.png
```

3. Installation

Create a Python virtual environment and install dependencies:

```
1 python -m venv .venv
2 source .venv/bin/activate # Windows: .venv\Scripts\activate
3 pip install --upgrade pip
4 pip install -r requirements.txt
```

The examples require NumPy, PyTorch, and Matplotlib. CUDA is optional; the checked-in runs are CPU-friendly.

4. Quick Smoke Test

Run the fast local check:

```
bash scripts/run_smoke_tests.sh
```

A smoke test checks syntax and basic data flow. It does not train high-accuracy PINNs or prove identifiability.

5. Normal Runs

5.1 Static figure generation

```
python scripts/generate_figures.py
```

Expected output: refreshed PNG figures under `figures/`, including sparse-data setup, PIDL missing-mechanism, and neural-architecture diagrams.

5.2 Inverse PINN

```
python src/inverse_pinn_sir_malware.py --iters 1000
```

Expected output: total loss, learned beta/gamma, data loss, ODE residual loss, and initial-condition loss.

5.3 PIDL missing mechanism

```
python src/pidl_unknown_mechanism.py --iters 1000
```

Expected output: total loss, data loss, residual loss, correction regularizer, and mean learned correction.

5.4 Direct control PINN

```
python src/control_pinn_malware.py --iters 1000
```

Expected output: objective proxy, state residual, initial-condition loss, and mean control intensity.

5.5 PMP-informed PINN

```
python src/pmp_informed_pinn_malware.py --iters 1000
```

Expected output: state residual, costate residual, stationarity residual, and boundary residual.

5.6 Longer training diagnostics

```
python scripts/run_training_iterations.py --iters 600
```

Expected output:

- CSV histories under `experiments/`;
- `OUTPUT_PREVIEW.md` and `training_summary.md`;
- `figures/training_iteration_diagnostics.png`.

6. How to Interpret Outputs

Output	Interpretation
Data loss	How well the learned state matches observed data. Low data loss alone does not prove the dynamics are right.
ODE residual loss	How well the neural derivative satisfies the differential equation at collocation points.
Initial-condition loss	Whether the learned trajectory starts from the prescribed cyber state.
Correction regularizer	Whether the PIDL correction remains controlled instead of replacing the known mechanism.
Objective proxy	The direct-control PINN's malware-control objective. It should be checked together with residual loss.
Stationarity loss	Whether the PMP-informed control satisfies the Hamiltonian stationarity condition.
Boundary loss	Whether the state initial condition and costate terminal condition are satisfied.

7. Output Map

Start with `experiments/OUTPUT_PREVIEW.md`. It groups the generated results by experiment, diagnostic panel, and file. Then inspect `figures/training_iteration_diagnostics.png` and the CSV file for the experiment you are studying.

File	What it shows
<code>inverse_pinn_training_history.csv</code>	Total, data, ODE, and initial-condition losses plus learned beta/gamma.
<code>pidl_training_history.csv</code>	Known-mechanism residual loss, correction regularization, and mean learned correction.
<code>control_pinn_training_history.csv</code>	Objective, state residual, initial-condition loss, and mean control.
<code>pmp_informed_pinn_training_history.csv</code>	State, costate, stationarity, and boundary residuals.

`training_iteration_` Four-panel comparison of the main loss and residual signals.
`diagnostics.png`

8. Extending the Code

8.1 From sparse inverse PINN to noisy cyber data

Add observation noise, hold out a validation set, and report parameter uncertainty. Avoid claiming exact parameter recovery unless the data are synthetic and the ground truth is known.

8.2 From PIDL correction to richer missing mechanisms

Change the correction network input from (t, x) to (t, x, u) or to graph features when the missing mechanism depends on interventions or network structure. Keep a correction regularizer so the learned term remains interpretable.

8.3 From open-loop to feedback neural control

Change the control network input from t to (t, x) , then train over multiple initial states and parameter scenarios. This produces a feedback-like neural control law rather than one open-loop control curve.

8.4 From small SIR models to network-scale PINNs

Move gradually: first degree-level compartments, then node-level vectors or graph-neural PINNs. For stiff, impulsive, or hybrid dynamics, use piecewise networks and add collocation points near jump times.

9. Troubleshooting

- If data loss is low but residual loss is high, increase residual weight or add collocation points.
- If residual loss is low but prediction is poor, the assumed model may be wrong or parameters may be unidentifiable.
- If PIDL correction dominates the known mechanism, increase correction regularization or reduce correction-network width.
- If direct-control PINN lowers the objective but violates dynamics, increase state-residual weight and validate the learned control with an independent ODE solver.
- If PMP-informed stationarity loss remains high, check control bounds and consider projected or KKT-style stationarity residuals.

10. Reproducibility Checklist

For every experiment, record:

- model equations and parameter values;
- random seed, initial states, and observation schedule;
- collocation-point sampling strategy;
- network architecture and optimizer;
- loss weights for data, residual, boundary, objective, and PMP terms;
- number of iterations and logging interval;
- all baseline solvers or held-out checks used for comparison.